




Results in Charts


Every chart below uses numbers measured live on one HP ZGX Nano (NVIDIA GB10) on Sep 25, 2026. Nano vs AWS figures come from 23 requests sent to both with identical prompts; AWS CloudWatch's own counts matched ours call for call. Public benchmark and synthetic data only.



Monthly cost at scale

Nano: \$3,999 box spread over 36 months + measured electricity · AWS: measured cost per request × volume

■ Nano ■ AWS on-demand ■ AWS enterprise*




Nano line stops at one Nano's measured capacity. Below break-even AWS is cheaper; above it the Nano is.

Jev vs the Nano check model, 13 live requests

Score each gave the request (0 = safe, 1 = attack) · Jev alone caught 5 attacks the Nano scored below 0.5


■ Jev (asked first) ■ Nano check model



Nano check model: precision and recall by dataset

10,632 held-back test requests · false-positive rate 0.09% overall

Precision Recall




* Path traversal was deliberately left out of training, so the learning loop has a real blind spot to close.

Tokens generated per second on the Nano

Both models serving at the same time (HP ZRT / vLLM), 50 requests per point

Qwen3-Next-80B writer Gemma 4 12B verifier



GPU memory in use (121.6 GB usable unified memory)

101.5 GB in use with both models serving · fine-tuning peaked at 20.4 GB




Fine-tuned patch writer on attack types it never saw


80 tasks · Qwen2.5-7B + LoRA trained on the Nano from 230 test-verified rules

■ Base 7B (baseline) ■ Fine-tuned 7B ■ 80B teacher

Rules passing every test
higher is better



Attack recall
higher is better



False positives
lower is better

